

Ques1:

Ans:

Mapper:

```
import java.io.IOException;
```

```
import java.util.StringTokenizer;
```

```
import org.apache.hadoop.io.DoubleWritable;
```

```
import org.apache.hadoop.io.IntWritable;
```

```
import org.apache.hadoop.io.LongWritable;
```

```
import org.apache.hadoop.io.Text;
```

```
import org.apache.hadoop.mapreduce.Mapper;
```

```
public class MyMapper extends Mapper<LongWritable, Text, Text, IntWritable>
```

```
{
```

```
    @Override
```

```
    protected void map(LongWritable key, Text value,
```

```
                        org.apache.hadoop.mapreduce.Mapper.Context context)
```

```
        throws IOException, InterruptedException {
```

```
        String inputstring = value.toString();
```

```
        for(String x : inputstring.split(" "))
```

```
{
```

```
            context.write(new Text(x),new IntWritable(1));
```

```
}
```

```
}
```

```
}
```

Reducer:

```
import java.io.IOException;
```

```
import org.apache.hadoop.io.IntWritable;
```

```
import org.apache.hadoop.io.Text;
```

```
import org.apache.hadoop.mapreduce.Reducer;
```

```
import org.apache.hadoop.mapreduce.Reducer.Context;
```

```
public class MyReducer extends Reducer<Text, IntWritable, Text, IntWritable>{
```

```
    @Override
```

```
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
```

```
        throws IOException, InterruptedException {
```

```
            int y = 0;
```

```
            for(IntWritable x : values)
```

```
            {
```

```
                y ++;
```

```
            }
```

```
            context.write(key, new IntWritable(y) );
```

```
        }
```

```
}
```

Driver:

```
import java.io.IOException;
```

```
import java.net.URI;
```

```
import java.net.URISyntaxException;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
//import org.apache.hadoop.mapred.FileInputFormat;
//import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class MyDriver {

    public static void main(String[] args) throws IOException, ClassNotFoundException,
    InterruptedException, URISyntaxException

    {

        Configuration conf = new Configuration();

        j.setJobName("My First Job");

        j.setJarByClass(MyDriver.class );

        j.setMapperClass(MyMapper.class );

        j.setReducerClass(MyReducer.class);

        j.setOutputKeyClass(Text.class);

        j.setOutputValueClass(IntWritable.class);

        FileInputFormat.addInputPath(j, new Path(args[0]));

        FileOutputFormat.setOutputPath(j, new Path(args[1]));
```

```
URI uri = new URI(args[1].toString());
```

```
FileSystem fs = FileSystem.get(uri, conf);
```

```
boolean x = fs.delete(new Path(uri),true);
```

```
int xxx = j.waitForCompletion(true) ? 0 : 1;
```

```
}
```

```
}
```

OUTPUT

The screenshot shows the Cloudera Hue File Browser interface. The browser is displaying the contents of a file named 'part-r-00000' located at the path '/output_WCdemo / part-r-00000'. The file is a text file containing a list of words and their frequencies. The interface includes a sidebar with actions like 'Edit file', 'Download', and 'View file location'. The main content area shows the text of the file, which is a list of words and their frequencies, such as 'All 1', 'Apache 2', 'Distributed 1', 'File 1', 'Hadoop 6', 'It 3', 'MapReduce 2', 'System 1', 'The 1', 'This 2', 'a 10', 'access 1', 'across 1', 'advantage 1', 'allows 1', 'also 1', 'amounts 1', 'and 7', 'approach 1', 'architecture 1', 'are 3', 'as 1', 'assumption 1', and 'automatically 1'.

File Browser

ACTIONS

- View as binary
- Edit file
- Download
- View file location
- Refresh

INFO

Last modified: June 18, 2021 9:43 p.m.

User: cloudera

Group: supergroup

Size: 1.2 KB

Mode: 100644

Home / output_WCdemo / part-r-00000

Page 1 of 1

```
( 1
(HDFS), 1
/ha'du:p/) 1
All 1
Apache 2
Distributed 1
File 1
Hadoop 6
It 3
MapReduce 2
System 1
The 1
This 2
a 10
access 1
across 1
advantage 1
allows 1
also 1
amounts 1
and 7
approach 1
architecture 1
are 3
as 1
assumption 1
automatically 1
```

Ques2-

//Driver

```
import java.io.IOException;
```

```
import org.apache.hadoop.conf.Configuration;
```

```
import org.apache.hadoop.fs.Path;
```

```
import org.apache.hadoop.io.IntWritable;
```

```
import org.apache.hadoop.io.Text;
```

```
import org.apache.hadoop.mapreduce.Job;
```

```
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
```

```
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

```
public class SumofEvenDriver
```

```
{
```

```
    public static void main(String[] args) throws IOException, ClassNotFoundException,  
    InterruptedException
```

```
    {
```

```
        //Path input_file_loc = new  
        Path("hdfs://localhost:9000/user/Vijay_Singh/evenodd.txt/"); //hardcoded input location
```

```
        //Path output_dir_loc = new  
        Path("hdfs://localhost:9000/user/Vijay_Singh/output/"); //hardcoded output location
```

```
        Configuration conf = new Configuration();
```

```
        Job job = new Job(conf, "Evenodd");
```

```
        ////name of Driver class
```

```
        job.setJarByClass(SumofEvenDriver.class);
```

```
        //name of mapper class
```

```
        job.setMapperClass(SumofEvenMapper.class);
```

```
        // name of reducer class
```

```

        job.setReducerClass(SumofEvenReducer.class);

        job.setMapOutputKeyClass(Text.class);

        job.setMapOutputValueClass(IntWritable.class);


        //FileInputFormat.addInputPath(job, input_file_loc);
        //FileOutputFormat.setOutputPath(job, output_dir_loc);
        //output_dir_loc.getFileSystem(job.getConfiguration()).delete(output_dir_loc,true);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));


        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

//Mapper
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class SumofEvenMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{

    public void map(LongWritable key, Text value, Context context)throws java.io.IOException,
    InterruptedException
    {

        String data[]=value.toString().split(",");  //data = [85 131 993 392 689....]


        for(String num:data)
        {

```

```

        int number=Integer.parseInt(num);
        if((number%2)==0)
        {
            context.write(new Text("EVEN"), new IntWritable(number)); //
ODD 85 131 993
        }

    }

}

//Reducer
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class SumofEvenReducer extends Reducer<Text, IntWritable, Text, IntWritable>
{

    // reducer will recieve key value pair as  ODD [ 85 131 993 491 539 985 413.....
    //      EVEN [ 392 870 240 888 184 494 704 996 408....

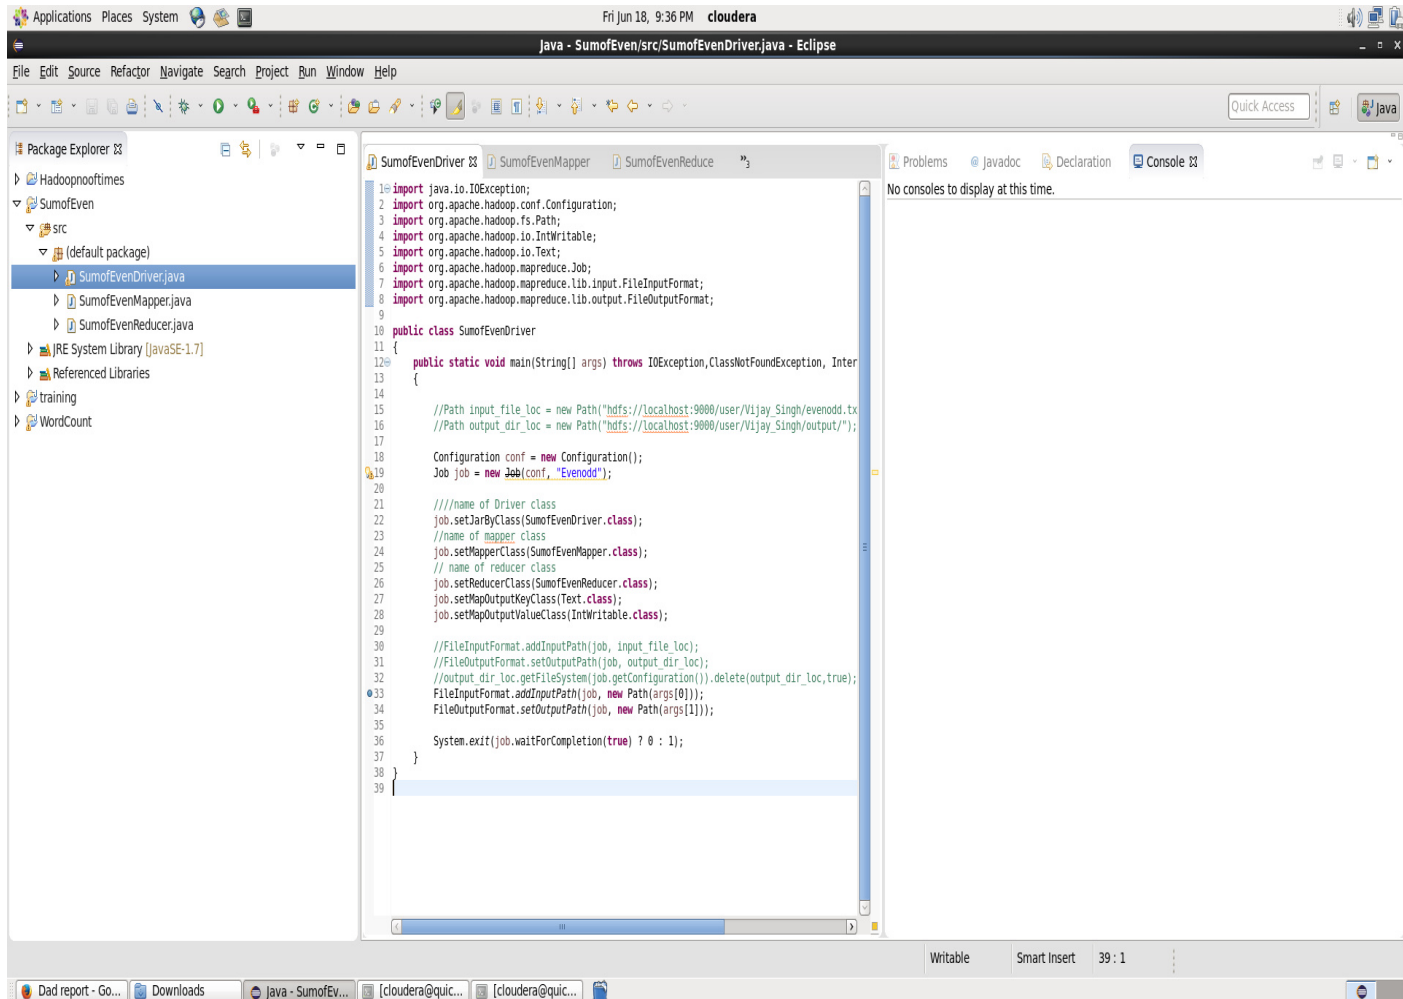
    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws
IOException, InterruptedException
    {
        int sum = 0;
        if(key.equals("EVEN"))
        {
            for (IntWritable value : values)
            {
                sum += value.get();
            }
        }
    }
}

```

```
context.write(key,new IntWritable(sum));;
```

```
}
```

```
}
```



Output

The screenshot displays the Hue File Browser interface. The top navigation bar includes links to Cloudera, Hue, Hadoop, HBase, Impala, Spark, Solr, Oozie, Cloudera Manager, and Getting Started. The main header shows the current path: Home / output_EDemo / part-r-00000. The sidebar on the left contains a list of actions (View as binary, Edit file, Download, View file location, Refresh) and file information (Last modified: June 18, 2021 9:18 p.m., User: cloudera, Group: supergroup, Size: 12 bytes, Mode: 100644). The main content area shows a table with the following data:

File Name	Size
part-r-00000	12 bytes

The bottom status bar shows the current session information: Hue - File Browser - pa..., Java - EvenOdd/src/My..., cloudera@quickstart..., and [Q1.txt (~Desktop) - g...].